

知识驱动大模型智能体 Knowledgeable Agents

大模型智能体


■ 大模型智能体研究正迎来推理规划技术落地的黄金时期

通过递归式目标分解与价值对齐机制，将模糊的人类需求转化为可执行的子任务



Deep Research能够主动拆解用户指令，引导用户确认真正意图，并利用搜索功能，基于互联网的实时信息做出回答。发现矛盾信息时能自主验证、回溯。

2025年2月2日，OpenAI发布新型智能体产品Deep Research。



Daily Papers

+ Submit a paper

by and the research community

Search any paper with AI

Daily

Weekly

Monthly

★

Sep 17

WebWeaver: Structuring Web-Scale Evidence with Dynamic Outlines for Open-Ended Deep Research

Zijun Li, Xin Guan, Bo Zhang, Shen Huang¹⁰⁰, Huiquan Zhou, Shaoping Lai, Ming Yan, Yong Jiang¹⁰⁰, Pengxin Xie, Fei Huang, Jun Zhang, Jingren Zhou

Tongyi Lab¹⁰⁰, Alibaba Group

<https://tongyi-agent.github.io/blog>
<https://github.com/Alibaba-ILP/DeepResearch>

Abstract

This paper tackles open-ended deep research (OEDR), a complex challenge where AI agents must synthesize vast web-scale information into insightful reports. Current approaches are plagued by dual-fold limitations: static research pipelines that decouple planning from evidence acquisition and...
Submitted by taesiri

35

WebWeaver: Structuring Web-Scale Evidence with Dynamic Outlines for Open-Ended Deep Research

12 authors

6.89k

1

Scaling Agents via Continual Pre-training

Liangcai Su¹⁰⁰, Zhen Zhang¹⁰⁰, Guangyu Li¹⁰⁰, Zhao Chen¹⁰⁰, Chenshi Wang¹⁰⁰, Maopu Song, Xinyu Wang¹⁰⁰, Kuan Li, Jialong Wu, Xuanzhong Chen, Zile Qiao, Zhongwang Zhang, Hailong Yin, Shihao Cai, Rannan Fang, Zhengwei Tao, Wenbiao Yin, Chenxiang Qian, Yong Jiang¹⁰⁰, Pengxin Xie, Fei Huang, Jingren Zhou

Tongyi Lab¹⁰⁰, Alibaba Group

<https://tongyi-agent.github.io/blog>
<https://github.com/Alibaba-ILP/DeepResearch>

Abstract

Large language models (LLMs) have evolved into agentic systems capable of autonomous tool use and multi-step reasoning for complex problem-solving. However, post-training approaches building upon general-purpose foundation models consistently underperform in agentic tasks, particularly in...
Submitted by callanwu

30

Scaling Agents via Continual Pre-training

22 authors

6.89k

2

WebSailor-V2: Bridging the Chasm to Proprietary Agents via Synthetic Data and Scalable Reinforcement Learning

Kuan Li¹⁰⁰, Zhongwang Zhang¹⁰⁰, Hailong Yin¹⁰⁰, Bai Yu¹⁰⁰, Yida Zhao¹⁰⁰, Liwen Zhang¹⁰⁰, Liu Ou, Dinghu Zhang, Xizi Wu, Jialong Wu, Xinyu Wang, Zile Qiao, Zhen Zhang, Yong Jiang¹⁰⁰, Pengxin Xie, Fei Huang, Jingren Zhou

Tongyi Lab¹⁰⁰, Alibaba Group

<https://tongyi-agent.github.io/blog>
<https://github.com/Alibaba-ILP/DeepResearch>

Abstract

To significantly advance the capabilities of open-source web agents, we present WebSailor-V2, a complete post-training pipeline encompassing data construction, Supervised Fine-Tuning (SFT), and Reinforcement Learning (RL). Our methodology features two key innovations: (1) On the data front, we developed Sailorgpt-QA-V2, a novel dataset built from a densely inter...
Submitted by callanwu

29

WebSailor-V2: Bridging the Chasm to Proprietary Agents via Synthetic Data and Scalable Reinforcement Learning

17 authors

6.89k

2

Towards General Agentic Intelligence via Environment Scaling

Rannan Fang¹⁰⁰, Shihao Cai¹⁰⁰, Baixuan Li¹⁰⁰, Jialong Wu¹⁰⁰, Guangyu Li¹⁰⁰, Wenbiao Yin, Xinyu Wang, Xiaobin Wang, Liangcai Su, Zhen Zhang, Shihao Wu, Zhengwei Tao, Yong Jiang¹⁰⁰, Pengxin Xie, Fei Huang, Jingren Zhou

Tongyi Lab¹⁰⁰, Alibaba Group

<https://tongyi-agent.github.io/blog>
<https://github.com/Alibaba-ILP/DeepResearch>

Abstract

Advanced agentic intelligence is a prerequisite for deploying Large Language Models in practical, real-world applications. Diverse real-world APIs demand precise, robust function-calling intelligence, which needs agents to develop these capabilities through interaction in varied environments...
Submitted by taesiri

26

Towards General Agentic Intelligence via Environment Scaling

16 authors

6.89k

2

WebResearcher: Unleashing unbounded reasoning capability in Long-Horizon Agents

Zile Qiao¹⁰⁰, Guankai Chen¹⁰⁰, Xuanzhong Chen¹⁰⁰, Dinghu Yu¹⁰⁰, Wenbiao Yin, Xinyu Wang, Zhen Zhang, Baixuan Li, Hailong Yin, Kuan Li, Bai Mu, Mingping Luo, Yong Jiang¹⁰⁰, Pengxin Xie, Fei Huang, Jingren Zhou

Tongyi Lab¹⁰⁰, Alibaba Group

<https://tongyi-agent.github.io/blog>
<https://github.com/Alibaba-ILP/DeepResearch>

Abstract

Recent advances in deep-research systems have demonstrated the potential for AI agents to autonomously discover and synthesize knowledge from external sources. In this paper, we introduce WebResearcher, a novel framework for building such agents through two key components: (1) An iterative deep-research paradigm that reformulates deep research...
Submitted by taesiri

26

WebResearcher: Unleashing unbounded reasoning capability in Long-Horizon Agents

16 authors

6.89k

2

ReSum: Unlocking Long-Horizon Search Intelligence via Context Summarization

Xizi Wu¹⁰⁰, Kuan Li¹⁰⁰, Yida Zhao, Liwen Zhang¹⁰⁰, Liu Ou, Hailong Yin, Zhongwang Zhang, Yong Jiang¹⁰⁰, Pengxin Xie, Fei Huang, Maobao Cheng, Shao Wang, Hong Cheng, Jingren Zhou

Tongyi Lab¹⁰⁰, Alibaba Group

<https://tongyi-agent.github.io/blog>
<https://github.com/Alibaba-ILP/DeepResearch>

Abstract

Large Language Model (LLM)-based web agents demonstrate strong performance on knowledge-intensive tasks but are hindered by context window limitations in paradigms like ReAct. Complex queries involving multiple entities, interrelated relationships, and high uncertainty demand extensive search cycles that rapidly exhaust context budgets before reaching complete solutions. To overcome this challenge, we introduce ReSum that enables indefinite exploration through periodic context summarization...
Submitted by callanwu

24

ReSum: Unlocking Long-Horizon Search Intelligence via Context Summarization

14 authors


6.89k

2

大模型智能体

- “假设大脑被取出置于营养液中，通过计算机模拟感官信号，使其无法区分虚拟与真实体验。”

希拉里·普特南，《理性，真理与历史》，1981年



缸中之脑



智能体

大模型智能体：应用之殇

真实应用场景的大模型智能体存在不稳定、不可靠等问题

ReAct

Question: Is Children's National Medical Center or MedStar Washington Hospital Center the largest private hospital in Washington, D.C.?

Thought 1:

I need to find who is the largest private hospital in Washington, D.C. by comparing Children's National Medical Center and MedStar Washington Hospital Center.

Action 1: Lookup[Children's National Medical Center] → Invalid

Observation 1: The last page Searched was not found, ...

Thought 2: Hmm, neither of the addresses is similar to any of the previous Searches. Maybe I need to rephrase the question.

Action 2: Search[question rephrased]


...

Action 3: Finish[not possible]

Observation 3: Answer is INCORRECT

Lookup --> Search ❌

Web/GUI智能体



大模型智能体：应用之殇

真实应用场景的大模型智能体存在不稳定、不可靠等问题

ReAct

Question: Is Children's National Medical Center or MedStar Washington Hospital Center the largest private hospital in Washington, D.C.?

Thought 1:

I need to find who is the largest private hospital in Washington, D.C. by comparing Children's National Medical Center and MedStar Washington Hospital Center.

Action 1: Lookup[Children's National Medical Center] → Invalid

Observation 1: The last page Searched was not found, ...

Thought 2: Hmm, neither of the addresses is similar to any of the previous Searches. Maybe I need to rephrase the question.

Action 2: Search[question rephrased]


...

Action 3: Finish[not possible]

Observation 3: Answer is INCORRECT

Lookup --> Search ❌

Web/GUI智能体



大模型智能体：应用之觞

真实应用场景的大模型智能体存在不稳定、不可靠等问题

Text

Image

Table

File

Question Inputs

DS Agent

Solution

Evaluation

An election has been held in Excelstan. Excelstan is a small country including 9 Districts. There are 1000 voters. Each voter is assigned a District Code based on where they live. The District Code is a number between 10000 and 10999 and determines what District the voter votes for.

No.	DistrictCode	District
1	105 ~ 114	Alpha
2	115 ~ 124	Beta
3	125 ~ 134	Gamma
4	135 ~ 144	Delta
5	145 ~ 154	Epsilon
6	155 ~ 164	Zeta
7	165 ~ 174	Eta
8	175 ~ 184	Theta
9	185 ~ 194	Iota

Table: districtcode in Excelstan. Excel file: Voter information and their voted district (not show).

How many voters are there in the Delta District?
A. 113 B. 114 C. 115 D. 116 E. 117

Given the problem description and data files, a DS agent generates executable codes to solve the problem.

```
jupyter Solution_for_Election_Voting Last Checkpoint: 6 hours ago
File Edit View Run Kernel Settings Help
import pandas as pd
import math

data_excel = pd.read_excel('Election_Voting.xlsx', sheet_name="Data")
data_excel

delta_voting_count = 0
for dis_code in data_excel['District Code']:
    if int(dis_code) >= 135 and int(dis_code) <= 144:
        delta_voting_count += 1
print('Number of voters in Delta district:', delta_voting_count)

Number of voters in Delta district: 115
1000 rows x 3 columns
```

After executing the code, the DS agent finalizes the answer based on the execution results: There are 115 voters in Delta district, the answer is C.

数据科学智能体

Models	Without Knowledge				With Knowledge			
	SR	CBS	VER	Cost ↓	SR	CBS	VER	Cost ↓
Direct Prompting								
Llama-3.1-Instruct-70B	5.9	81.5	29.4	0.001	4.9	82.1	27.5	0.001
Llama-3.1-Instruct-405B	3.9	79.4	35.3	0.010	2.9	81.3	25.5	0.011
Mistral-Large-2 (2407)	13.7	83.2	47.1	0.009	16.7	84.7	39.2	0.009
GPT-4o (2024-05-13)	11.8	82.6	52.9	0.011	10.8	83.8	41.2	0.012
Claude-3.5-Sonnet (2024-06-20)	17.7	83.6	51.0	0.017	21.6	85.4	41.2	0.017
OpenAI o1†	34.3	87.1	70.6	0.221	31.4	87.4	63.7	0.236
OpenHands CodeAct								
Llama-3.1-Instruct-70B	6.9	63.5	30.4	0.145	2.9	65.7	25.5	0.252
Llama-3.1-Instruct-405B	5.9	65.8	52.0	0.383	8.8	71.4	58.8	0.740
Mistral-Large-2 (2407)	9.8	72.5	53.9	0.513	13.7	78.8	50.0	0.759
GPT-4o (2024-05-13)	19.6	83.1	78.4	0.803	27.5	86.3	73.5	1.094
Claude-3.5-Sonnet (2024-06-20)	21.6	83.6	87.3	0.958	24.5	85.1	88.2	0.900
Self-Debug								
Llama-3.1-Instruct-70B	13.7	82.7	80.4	0.007	16.7	83.4	73.5	0.008
Llama-3.1-Instruct-405B	14.7	82.9	78.4	0.047	13.7	83.6	79.4	0.055
Mistral-Large-2 (2407)	23.5	85.1	83.3	0.034	27.5	86.8	78.4	0.036
GPT-4o (2024-05-13)	22.6	84.4	83.3	0.047	23.5	85.6	71.6	0.046
Claude-3.5-Sonnet (2024-06-20)	32.4	86.4	92.2	0.057	34.3	87.1	86.3	0.061
OpenAI o1†	42.2	88.4	92.2	0.636	41.2	88.9	91.2	0.713

- 1 DS Bench: How Far Are Data Science Agents to Becoming Data Science Experts? (ICLR 2025)
- 2 ScienceAgentBench: Toward rigorous assessment of language agents for data-driven scientific discovery (ICLR 2025)

大模型智能体：解决之路在何方？

workflows? 强化学习?

**Cognitive Behaviors that Enable Self-Improving Reasoners,
or,
Four Habits of Highly Effective STaRs**

Kanishk Gandhi
Stanford University

Ayush Chakravarthy
Stanford University

Anikait Singh
Stanford University

Nathan Lile
SynthLabs

Noah D. Goodman
Stanford University

**Does Reinforcement Learning Really Incentivize Reasoning
Capacity in LLMs Beyond the Base Model?**

Yang Yue^{1*†}, Zhiqi Chen^{1*}, Rui Lu¹, Andrew Zhao¹, Zhaokai Wang², Yang Yue¹,
Shiji Song¹, and Gao Huang^{1✉}

¹ LeapLab, Tsinghua University ² Shanghai Jiao Tong University

* Equal Contribution † Project Lead ✉ Corresponding Author

Engineering at Anthropic



Building effective agents

OpenAI



**A practical
guide to
building agents**

知识增强智能体的必要性

为什么AI智能体推理规划需要知识？


智能体产生不可靠决策的根本原因在于**统计模型在不确定性预测上的固有限制以及先验知识的缺失**

❑ 长链推理中出现错误（计算、语法、格式等）、反思能力不足难以自纠



知识是行业领域智能体应用精准性和可靠性的重要基础和基本保障

知识与智能体的历史渊源



知识驱动智能体的技术路线

知识驱动智能体的基本范式

要让大模型更好的“懂得”行业知识，需要更为精细的**知识增强**方法

□ **基础知识增强**: 构造数据训练（见多识广）

□ **输入先验增强**: 提示增强（有的放矢）

□ **输出生成约束**: 解码约束（循规蹈矩）


数据构建



训练推理



记忆（知识）更新




符号知识




参数知识（大模型）

通过数据驱动或奖励塑造实现知识获取，而非人工构建，支撑行业智能体应用场景


知识驱动智能体的技术栈




知识驱动智能体的应用场景



(a) Overview



(b) Single Agent



(c) Multi-Agents

知识驱动智能体

知识驱动
智能体

```
graph LR; A[知识驱动智能体] --> B[推理与规划]; A --> C[干预与更新]; A --> D[创新与创造];
```

The diagram illustrates the components of a knowledge-driven intelligent agent. A central grey box labeled '知识驱动智能体' (Knowledge-Driven Intelligent Agent) has three arrows pointing to three red boxes: '推理与规划' (Reasoning and Planning), '干预与更新' (Intervention and Update), and '创新与创造' (Innovation and Creation). Each red box is followed by a grey box containing a descriptive text.

推理与规划

如何实现知识增强推理与可靠规划

干预与更新

如何实现知识引导行为干预与更新

创新与创造

如何实现知识创新与新方法的创造

知识驱动智能体

知识驱动
智能体

```
graph LR; A[知识驱动智能体] --> B[推理与规划]; A --> C[干预与更新]; A --> D[创新与创造];
```

The diagram illustrates the components of a knowledge-driven intelligent agent. A central grey box labeled '知识驱动智能体' (Knowledge-driven Intelligent Agent) has three arrows pointing to three separate horizontal bars. Each bar consists of a red box on the left and a light grey box on the right. The top bar is highlighted with a red border. The three bars represent the core functions: reasoning and planning, intervention and update, and innovation and creation.

推理与规划

如何实现知识增强推理与可靠规划

干预与更新



如何实现知识引导行为干预与更新

创新与创造


如何实现知识创新与新方法的创造

知识驱动智能体：通过构建数据训练增强知识

智能体思维链与轨迹的合成与筛选



轨迹获取与合成




流程知识抽取


- 1 PC Agent: While You Sleep, AI Works - A Cognitive Journey into Digital World
- 2 Annotation and Extraction of Industrial Procedural Knowledge from Textual Documents

KnowAgent

知识增强智能体促进可靠决策：构造数据训练、提示增强




知识增强智能体促进可靠决策：构造数据训练、提示增强




- ❑ Convert task-specific action knowledge into text descriptions and use this knowledge to guide the model in synthesizing planning paths (**Planning Path Generation**).
- ❑ The model then iteratively optimizes its performance using the generated planning trajectories (**Knowledgeable Self-Learning**).

知识增强智能体促进可靠决策：构造数据训练、提示增强



- ❑ We convert task-specific action knowledge into text descriptions and use this knowledge to guide the model in synthesizing planning paths (**Planning Path Generation**).
- ❑ The model then iteratively optimizes its performance using the generated planning trajectories (**Knowledgeable Self-Learning**).

知识增强智能体促进可靠决策：构造数据训练、提示增强



- ❑ We convert task-specific action knowledge into text descriptions and use this knowledge to guide the model in synthesizing planning paths (**Planning Path Generation**).
- ❑ The model then iteratively optimizes its performance using the generated planning trajectories (**Knowledgeable Self-Learning**).

KnowAgent

Backbone	Strategy	Method	HotpotQA				ALFWorld	
			Easy	Medium	Hard	Average	Seen	Unseen
Llama-2 7b-chat	Prompting	CoT (Wei et al., 2022)	35.80	26.69	18.20	26.90	-	-
	Prompting	ReAcT (Yao et al., 2023b)	25.14	19.87	17.39	20.80	1.43	0.75
	Prompting	Reflexion (Shinn et al., 2023)	35.55	28.73	24.35	29.54	2.86	2.24
	Fine-tuning	FireAct (Chen et al., 2023a)	40.56	31.70	24.13	32.13	27.86	25.38
	Fine-tuning	KNOWAGENT-7B	40.80	32.49	27.12	33.47	29.26	29.35
Llama-2 13b-chat	Prompting	CoT (Wei et al., 2022)	37.90	25.28	21.64	28.27	-	-
	Prompting	ReAcT (Yao et al., 2023b)	28.68	22.15	21.69	24.17	13.58	9.70
	Prompting	Reflexion (Shinn et al., 2023)	44.43	37.50	28.17	36.70	18.57	17.16
	Fine-tuning	FireAct (Chen et al., 2023a)	51.95	33.93	28.88	38.26	47.86	50.37
	Fine-tuning	KNOWAGENT-13B	46.97	37.60	33.22	39.26	54.29	58.71
Llama-2 70b-chat	Prompting	CoT (Wei et al., 2022)	45.37	36.33	32.27	37.99	-	-
	Prompting	ReAcT (Yao et al., 2023b)	39.70	37.19	33.62	36.83	41.43	37.31
	Prompting	Reflexion (Shinn et al., 2023)	48.01	46.35	35.64	43.33	49.29	47.01
	Fine-tuning	FireAct (Chen et al., 2023a)	51.96	47.56	44.60	48.04	77.85	77.61
	Fine-tuning	KNOWAGENT-70B	56.75	49.90	37.76	48.14	77.14	78.36

Table 1: **Overall performance of KNOWAGENT on HotpotQA and ALFWorld.** The evaluation metrics are F1 Score (%) and Success Rate (%), respectively. **Strategy** means the agent learning paradigm behind each method. The best results of each backbone are marked in **bold**.

KnowAgent demonstrates certain advantages comparing to fine-tuning methods

KnowAgent

Backbone	Strategy	Method	HotpotQA				ALFWorld	
			Easy	Medium	Hard	Average	Seen	Unseen
Llama-2 7b-chat	Prompting	CoT (Wei et al., 2022)	35.80	26.69	18.20	26.90	-	-
	Prompting	ReAcT (Yao et al., 2023b)	25.14	19.87	17.39	20.80	1.43	0.75
	Prompting	Reflexion (Shinn et al., 2023)	35.55	28.73	24.35	29.54	2.86	2.24
	Fine-tuning	FireAct (Chen et al., 2023a)	40.56	31.70	24.13	32.13	27.86	25.38
	Fine-tuning	KNOWAGENT-7B	40.80	32.49	27.12	33.47	29.26	29.35
Llama-2 13b-chat	Prompting	CoT (Wei et al., 2022)	37.90	25.28	21.64	28.27	-	-
	Prompting	ReAcT (Yao et al., 2023b)	28.68	22.15	21.69	24.17	13.58	9.70
	Prompting	Reflexion (Shinn et al., 2023)	44.43	37.50	28.17	36.70	18.57	17.16
	Fine-tuning	FireAct (Chen et al., 2023a)	51.95	33.93	28.88	38.26	47.86	50.37
	Fine-tuning	KNOWAGENT-13B	46.97	37.60	33.22	39.26	54.29	58.71
Llama-2 70b-chat	Prompting	CoT (Wei et al., 2022)	45.37	36.33	32.27	37.99	-	-
	Prompting	ReAcT (Yao et al., 2023b)	39.70	37.19	33.62	36.83	41.43	37.31
	Prompting	Reflexion (Shinn et al., 2023)	48.01	46.35	35.64	43.33	49.29	47.01
	Fine-tuning	FireAct (Chen et al., 2023a)	51.96	47.56	44.60	48.04	77.85	77.61
	Fine-tuning	KNOWAGENT-70B	56.75	49.90	37.76	48.14	77.14	78.36

Table 1: **Overall performance of KNOWAGENT on HotpotQA and ALFWorld.** The evaluation metrics are F1 Score (%) and Success Rate (%), respectively. **Strategy** means the agent learning paradigm behind each method. The best results of each backbone are marked in **bold**.

KnowAgent demonstrates certain advantages comparing to fine-tuning methods

KnowAgent

Method	ALFWorld	
	<i>Seen</i>	<i>Unseen</i>
ReAct	13.58	9.70
KNOWAGENT*	30.00 $\uparrow 16.42$	38.06 $\uparrow 28.36$

Table 2: Comparative Experiment of ReAct and KNOWAGENT* on ALFWorld with Llama-2-13b. Here label * denotes the version of KNOWAGENT integrated with action knowledge but not fine-tuned.

❑ Unreasonable Action Rates

Model	Invalid Action	Misordered Action
ReAct	2.08%	3.54%
Reflexion	6.87%	3.87%
KNOWAGENT	0.35%	1.23%

Table 3: Unreasonable action rates on HotpotQA with Llama-2-13b. Here *invalid* refers to actions that do not meet the action rule, while *misordered* means discrepancies in the logical sequence of actions.

The results validate the effectiveness of action knowledge itself.

❑ Distilled Knowledge vs. Manually Designed Knowledge

KNOWAGENT	Strategy	HotpotQA	ALFWorld
7B	Manual	33.47	20.15
	Distilled	25.22	18.66
13B	Manual	39.26	52.24
	Distilled	25.72	51.49


Table 4: Comparative Experiment on Manual vs. Distilled Action Knowledge. *Manual* stands for human-crafted knowledge and *Distilled* represents the distilled knowledge from GPT-4.

Action knowledge effectively mitigates planning hallucinations.

Utilizing advanced LLMs can supplant manual efforts in constructing task-specific action knowledge.


知识增强智能体促进新环境的泛化：脑补环境、优化动作/决策知识

Unfamiliar
Environment



Inadequate
Action Knowledge

知识增强智能体促进新环境的泛化：脑补环境、优化动作/决策知识



➤ Virtual scenarios Synthesis

➤ Multi-Tool use

➤ Real Environment FeedBack.

➤ MCTS-Refine.

Figure 2: **The overall framework of SynWorld:** we first extract composable tools from the toolkit to generate new scenes and tasks. Then, we allow agents to explore the synthesized virtual scenes using MCTS to optimize action knowledge, thereby learning how to execute actions and plan tasks.

知识增强智能体促进新环境的泛化：脑补环境、优化动作/决策知识

Action Knowledge:



Virtual Scenario Synthesis

$$\mathcal{S}(t) = \{(\mathcal{B}, \mathcal{G}) \mid \forall t \subseteq T\},$$

Background B: The contextual scenario specifying initial conditions and constraints;

Goal G: The terminal objective requiring tool-mediated resolution

Similarity Detection

$$d((\mathcal{B}_i, \mathcal{G}_i), (\mathcal{B}_j, \mathcal{G}_j)) < \epsilon.$$

知识增强智能体促进新环境的泛化：脑补环境、优化动作/决策知识

Initialization:

root node is initialized with predefined Action Knowledge


Feedback Collection:

$$Tra_i, S_i = Env(\mathcal{AK}, \pi)$$

$$\mathcal{E} = \{(S_{\text{before}}^i, S_{\text{after}}^i, \mathcal{M}^i) \mid N_i \in \text{Path}(N, N_0)\}$$

Exploration:

$$\mathcal{AK}_{\text{new}} = \pi_{\theta}(\mathcal{AK}_{\text{old}}, \mathcal{E}, Tra)$$



知识增强智能体促进新环境的泛化：脑补环境、优化动作/决策知识

Model	Method	ToolBench		HotpotQA
		PASS	WIN	
GPT-4-turbo	ReAct	50.67	67.00	54.61
	Self-Refine	56.80	73.00	55.85
	EasyTool	51.67	68.00	58.19
	DRAFT	54.83	72.00	57.71
	Ours	59.33	73.00	59.93
Qwen-long	ReAct	48.30	71.00	52.00
	Self-Refine	53.70	77.00	56.10
	EasyTool	50.80	63.00	58.34
	DRAFT	54.20	79.00	53.23
	Ours	57.20	81.00	59.91
Qwen2-72B-Instruct	ReAct	49.43	55.00	50.21
	Self-Refine	54.33	65.00	52.59
	EasyTool	52.97	58.00	54.94
	DRAFT	56.43	69.00	57.57
	Ours	58.52	73.00	58.70

Table 1: Main results of SynWorld compared to other

Model	Method	Pass Rate
GPT-4-turbo	SynWorld	59.33
	w/o. Workflow	56.33-3.00
	w/o. Description	53.16-6.17
Qwen-long	SynWorld	57.20
	w/o. Workflow	57.00-0.20
	w/o. Description	53.83-3.37

Table 2: Ablation experiment results

[1] SynWorld: Virtual Scenario Synthesis for Agentic Action Knowledge Refinement (ACL 2025)

知识增强智能体促进新环境的泛化：脑补环境、优化动作/决策知识




Figure 3: The variation in the pass rate of agents on the ToolBench in relation to the number of exploration scenarios.




Figure 4: Changes in ToolBench pass rates in virtual and real-world scenarios with the number of iterative optimizations performed in the virtual environment.

自主切换快慢思维 (KnowSelf)

Mehrdad Farajtabar
@MFarajtabar

1/ Can Large Language Models (LLMs) truly reason? Or are they just sophisticated pattern matchers? In our latest preprint, we explore this key question through a large-scale study of both open-source like Llama, Phi, Gemma, and Mistral and leading closed models, including the recent OpenAI GPT-4o and o1-series.
arxiv.org/pdf/2410.05229

Work done with @i_mirzadeh, @KeivanAlizadeh2, Hooman Shahrokhi, Samy Bengio, @OncelTuzel.


#LLM #Reasoning #Mathematics #AGI #Research #Apple

1/ 大型语言模型 (LLMs) 真的可以推理吗? 或者它们只是复杂的模式匹配器? 在我们最新的预印本中, 我们通过对 Llama、Phi、Gemma 和 Mistral 等开源模型以及领先的封闭模型 (包括最近的 OpenAI GPT-4o 和 o1 系列) 的大规模研究来探索这个关键问题。

完成的工作 @i_mirzadeh, @KeivanAlizadeh2, 胡曼·沙罗基, 萨米·本吉奥, @OncelTuzel。


#LLM #推理 #数学 #AGI #研究 #苹果

公众号·新智元



Transformer不会推理, 只是高级模式匹配器!

无脑的喂知识智能体表现反而变差!




Agentic Knowledgeable
Self-awareness

自主切换快慢思维 (KnowSelf)

提升智能体的“知识边界感知”能力

Step 1

Self-awareness Data Construction




Step 2

Self-awareness Learning



Step 3

Self-awareness Inference



Firstly, we mark self-explored trajectories with special tokens according to the situation judgement criterion to form the training data. **Secondly**, we apply a two-stage training framework to teach the agent model knowledgeable self-awareness abilities.

自主切换快慢思维 (KnowSelf)

提升智能体的 “知识边界感知” 能力

Backbone	Method	Know %	Put	Clean	Heat	Cool	Examine	Put Two	All
GPT-4o	✳ REACT	0%	83.33	74.19	69.57	85.71	77.78	64.71	76.12
	✳ Reflexion	0%	100.00	87.10	73.91	90.48	83.33	70.59	85.07
	✳ ExpeL	100%	95.83	83.87	69.57	80.95	88.89	52.94	79.85
Llama-8B	✳ REACT	0%	33.33	3.23	0.00	57.14	66.67	23.53	27.61
	✳ Reflexion	0%	62.96	22.73	5.88	64.29	86.36	50.00	51.49
	✳ ExpeL	100%	83.33	32.26	30.43	23.81	55.56	17.65	41.04
	💧 ETO	0%	91.67	70.59	82.61	61.90	88.89	64.71	78.36
	💧 KnowAgent	100%	87.50	93.55	65.22	66.67	61.11	64.71	75.37
	💧 WKM	100%	95.83	87.10	86.96	61.90	66.67	52.94	77.61
	💧 KnowSelf	15.01%	91.67	87.10	91.30	85.71	77.78	64.71	84.33
Gemma-2B	✳ REACT	0%	0.00	9.68	0.00	4.76	44.44	0.00	8.96
	✳ Reflexion	0%	4.76	10.71	0.00	9.52	65.38	0.00	17.16
	✳ ExpeL	100%	0.00	3.23	0.00	0.00	27.78	0.00	4.48
	💧 ETO	0%	91.67	83.87	78.26	52.38	77.78	29.41	71.64
	💧 KnowAgent	100%	91.67	90.32	69.57	71.43	66.67	41.18	73.88
	💧 WKM	100%	91.67	87.10	78.26	71.43	61.11	52.94	76.12
	💧 KnowSelf	26.41%	87.50	93.55	73.91	76.19	83.33	52.94	79.85

[1] Agentic Knowledgeable Self-awareness (ACL 2025)

抑制幻觉

提升智能体的“知识边界感知”能力

Why Language Models Hallucinate

Adam Tauman Kalai*
OpenAI


Ofir Nachum
OpenAI

Santosh S. Vempala†
Georgia Tech


Edwin Zhang
OpenAI

September 4, 2025

(a) Slow thinking model: **High reasoning**; **High hallucination**



(b) KnowRL (ours): **High reasoning**; **Low hallucination**



Learning to Reason for Factuality

Xilun Chen¹, Ilia Kulikov¹, Vincent-Pierre Berges¹, Barlas Oğuz¹, Rulin Shao², Gargi Ghosh¹, Jason Weston¹, Wen-tau Yih¹

¹FAIR at Meta, ²University of Washington

Reasoning Large Language Models (R-LLMs) have significantly advanced complex reasoning tasks but often struggle with factuality, generating substantially more hallucinations than their non-reasoning counterparts on long-form factuality benchmarks. However, extending online Reinforcement Learning (RL), a key component in recent R-LLM advancements, to the long-form factuality setting poses several unique challenges due to the lack of reliable verification methods. Previous work has utilized automatic factuality evaluation frameworks such as FActScore to curate preference data in the offline RL setting, yet we find that directly leveraging such methods as the reward in online RL leads to reward hacking in multiple ways, such as producing less *detailed* or *relevant* responses. We propose a novel reward function that simultaneously considers the factual precision, response detail level, and answer relevance, and applies online RL to learn high quality factual reasoning. Evaluated on six long-form factuality benchmarks, our factual reasoning model achieves an average reduction of **23.1 percentage points** in hallucination rate, a **23%** increase in answer detail level, and no degradation in the overall response helpfulness.

Date: August 8, 2025

Correspondence: Xilun Chen xilun@meta.com

Meta

- 1 KnowRL: Exploring Knowledgeable Reinforcement Learning for Factuality
- 2 Learning to Reason for Factuality (Meta, 2025)
- 3 Why Language Models Hallucinate (OpenAI)

抑制幻觉

提升智能体的“知识边界感知”能力

Methods	Hallucination			Reasoning	
	TruthfulQA	SimpleQA	ChineseSimpleQA	GPQA Diamond	AIME 2025
Skywork-OR1-7B-Preview					
Direct Generation	34.39	5.19	13.27	31.31	6.67
Self-Refine	30.00	2.90	11.57	46.97	36.67
FactTune-FS	34.76	5.33	12.30	32.83	6.67
DPO	34.64	4.60	15.57	34.34	6.67
SFT	38.43	4.90	14.87	32.32	16.67
KnowRL-Zero	34.39	5.55	15.17	30.30	10.00
KnowRL	37.70	4.58	16.23	34.34↑	10.00↑
DeepSeek-R1-Distill-Qwen-7B					
Direct Generation	33.05	1.92	11.17	34.34	16.67
Self-Refine	33.90	2.60	9.97	45.45	33.33
FactTune-FS	35.50	2.46	10.27	36.87	26.67
DPO	36.96	3.21	12.33	37.37	13.33
SFT	37.09	4.12	12.70	30.30	16.67
KnowRL-Zero	33.54	1.34	11.00	34.34	6.67
KnowRL	38.19	4.56	13.90	37.37↑	13.33↓

[1] KnowRL: Exploring Knowledgeable Reinforcement Learning for Factuality

知识驱动智能体: 知识增强 (强化) 学习

提升智能体的 “知识边界感知” 能力的终极情况: 意识?

Answer the given question. You must conduct reasoning inside `<think>` and `</think>` first every time you get new information. After reasoning, if you find you lack some knowledge, you can call a search engine by `<search>` query `</search>`, and it will return the top searched results between `<information>` and `</information>`. You can search as many times as you want. If you find no further external knowledge needed, you can directly provide the answer inside `<answer>` and `</answer>` without detailed illustrations. For example, `<answer>` xxx `</answer>`. Question: **question**.

Search-R1

System Prompt Template For Instruction-Tuned Model

You are a helpful assistant that can solve the given question step by step with the help of the wikipedia search tool. Given a question, you need to first think about the reasoning process in the mind and then provide the answer. During thinking, you can invoke the wikipedia search tool to search for fact information about specific topics if needed. The reasoning process and answer are enclosed within `<think>` `</think>` and `<answer>` `</answer>` tags respectively, and the search query and result are enclosed within `<search>` `</search>` and `<result>` `</result>` tags respectively. For example, `<think>` This is the reasoning process. `</think>` `<search>` search query here `</search>` `<result>` search result here `</result>` `<think>` This is the reasoning process. `</think>` `<answer>` The final answer is `\boxed{answer here}` `</answer>`. In the last part of the answer, the final exact answer is enclosed within `\boxed{}` with latex format.

ReSearch



- 1 Search-R1: Training LLMs to Reason and Leverage Search Engines with Reinforcement Learning
- 2 ReSearch: Learning to Reason with Search for LLMs via Reinforcement Learning

知识驱动智能体

知识驱动
智能体

```
graph LR; A[知识驱动智能体] --> B[推理与规划]; A --> C[干预与更新]; A --> D[创新与创造];
```

The diagram illustrates the components of a knowledge-driven intelligent agent. A central grey rounded rectangle labeled '知识驱动智能体' (Knowledge-driven Intelligent Agent) has three arrows pointing to three separate horizontal bars. Each bar consists of a red box on the left and a light grey box on the right. The top bar is '推理与规划' (Reasoning and Planning) with the description '如何实现知识增强推理与可靠规划' (How to achieve knowledge-enhanced reasoning and reliable planning). The middle bar is '干预与更新' (Intervention and Update), which is highlighted with a red border, with the description '如何实现知识引导行为干预与更新' (How to achieve knowledge-guided behavioral intervention and update). The bottom bar is '创新与创造' (Innovation and Creation) with the description '如何实现知识创新与新方法的创造' (How to achieve knowledge innovation and the creation of new methods).

推理与规划

如何实现知识增强推理与可靠规划

干预与更新



如何实现知识引导行为干预与更新

创新与创造

如何实现知识创新与新方法的创造

记忆（知识）更新

记忆动态更新、知识持续学习、知识编辑




智能体可靠规划的机理是什么？智能体如何进行知识持续学习和更新其能力？

- 1 Knowledge Circuits in Pretrained Transformers (NeurIPS 2024)
- 2 Editing Large Language Models: Problems, Methods, and Opportunities (EMNLP 2023)
- 3 A Comprehensive Study of Knowledge Editing for Large Language Models (2024)
- 4 The rise and potential of large language model based agents: A survey (Science China Information Sciences 2025)

知识编辑

知识编辑问题定义


- 通过对**符号**或**参数**知识的**新增、修改和擦除**等操作解决知识谬误问题，实现**可信、可控、可靠**的应用
- **三种模式**： I.全新知识-新增模式 II.有害知识-擦除模式 III.错误知识-修改模式



大模型知识编辑



难更新!



Knowledge Editing Types: Insertion Modification Erasure




<https://github.com/zjunlp/EasyEdit>

部署后的大模型存在知识截止、谬误、幻觉等一系列问题，知识编辑旨在高效、精准地更新（新增、擦除）大语言模型中的知识

知识编辑: 一种大语言模型记忆 (Memory) 更新技术

□ 适合语言模型的知识更新机制是什么?

认知科学: 人类记忆结构



1. 从工作记忆中习得长期记忆
2. 从长期记忆中检索工作记忆



LLM: 迈向可自我编辑更新的记忆结构


Working memory and language: An overview (Journal of communication disorders 2003)

WISE: Rethinking the Knowledge Memory for Lifelong Model Editing of Large Language Models (NeurIPS 2024)

持续学习：知识记忆


□ 已有的大模型记忆系统: **知识关联弱、更新效率低**

□ 最近的技术路线: 通过神经符号**知识编辑**驱动智能体, 实现新环境适应与持续学习




MEMOS: A Memory OS for AI System

Zhiyu Li^{1,3}, Shichao Song^{1,8}, Chenyang Xi¹, Hanyu Wang^{1,8}, Chen Tang¹, Simin Niu^{1,8}, Ding Chen¹⁰, Jiawei Yang^{1,8}, Chunyu Li¹, Qingchen Yu^{1,9}, Jihao Zhao^{1,8}, Yezhaohui Wang¹, Peng Liu⁸, Zehao Lin^{1,3}, Pengyuan Wang¹, Jiahao Huo¹, Tianyi Chen^{1,2}, Kai Chen^{1,3}, Kehang Li^{1,2}, Zhen Tao⁸, Junpeng Ren¹, Huayi Lai¹, Hao Wu¹, Bo Tang¹, Zhengren Wang^{7,3}, Zhaoxin Fan⁹, Ningyu Zhang⁵, Linfeng Zhang², Junchi Yan², Mingchuan Yang¹⁰, Tong Xu⁶, Wei Xu⁸, Huajun Chen⁵, Haofeng Wang⁴, Hongkang Yang^{1,3}, Wentao Zhang^{7,†}, Zhi-Qin John Xu^{2,†}, Siheng Chen^{2,†}, Feiyu Xiong^{1,3,†}



干预: 大模型行为调控与编辑方法




步骤1: 获取**调控向量**

步骤2: 选择**干预位置**


步骤3: 调控**模型行为**

大模型行为调控与编辑: 技术价值

□ 抽象知识（安全、人格）：“扭动旋钮” 运行时调控模型生成内容



Prompt?
Prompting is labor-intensive and sensitive




Steering is **interpretable, robust, and flexible**

大模型行为调控与编辑无需重训模型参数，具备实时精准解决问题的能力

大模型行为调控与编辑: 科学原理

挑战1: 如何找表征?
挑战2: 如何调行为?


抽象知识 (安全、人格): “扭动旋钮” 运行时调控模型生成内容



基于表征空间线性可操控性的科学假设, 揭示模型行为的可解释机理与可干预规律

大模型行为定向编辑干预方法: STA

SAE: disentangle




仅适用于相对简单的任务:

- 动词时态
- 实体识别
- 概念识别

大模型行为定向编辑干预方法: STA

□ 比较正负样本的神经元激活差异并设定阈值筛选，**锁定目标行为关键神经元**

1 定位



“原子” 激活**幅度**

$$\Delta \mathbf{a} = \frac{1}{N} \sum_{i=1}^N (\bar{\mathbf{a}}_{\text{pos}}^i - \bar{\mathbf{a}}_{\text{neg}}^i)$$

$$\mathbf{f}_j^{\text{pos}} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(|\bar{\mathbf{a}}_{j,\text{pos}}^i| > 0)$$

$$\mathbf{f}_j^{\text{neg}} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(|\bar{\mathbf{a}}_{j,\text{neg}}^i| > 0)$$

$$\Delta \mathbf{f} = \mathbf{f}_{\text{pos}} - \mathbf{f}_{\text{neg}}$$

“原子” 方向**频率**

2 干预

$$\mathbf{v}_{\text{STA}} = \mathbf{a}_{\text{target}} \mathbf{W}_{\text{dec}} + \mathbf{b}_{\text{dec}},$$

$$\hat{\mathbf{h}} = \lambda \mathbf{v}_{\text{STA}} + \mathbf{h}, \quad y = \mathcal{M}(x, \hat{\mathbf{h}}),$$



创新点：将稀疏表示与神经元激活残差相结合，从“原子粒度”上**解耦和定向模型内部表征**，从而在开放生成任务中实现可控且性能保持的行为干预

大模型行为定向编辑干预方法: STA

Model	Method	Detoxification Performance (↑)			General Performance (↑)			
		SafeEdit	RealToxicprompts	Avg	Fluency	MMLU	GSM8K	Avg
Gemma-2-9b-pt	Vanilla	62.30	57.63	59.97	4.31	62.34	67.55	44.73
	Prompt _{hand}	72.52	53.96	63.24	3.88	57.01	67.48	42.79
	Prompt _{auto}	64.15	57.63	60.89	4.19	60.09	68.61	44.30
	CAA	85.78	73.98	79.88	4.38	61.35	68.54	44.76
	SAE _{AXBENCH}	<u>86.81</u>	<u>75.15</u>	<u>80.98</u>	<u>4.33</u>	62.60	69.07	45.33
	STA (Ours)	89.93	76.98	83.45	4.29	<u>62.35</u>	65.05	43.90
Gemma-2-9b-it	Vanilla	70.37	97.41	83.89	5.39	72.06	75.66	51.04
	Prompt _{hand}	78.74	98.42	88.58	5.41	71.07	74.83	50.44
	Prompt _{auto}	75.56	<u>98.92</u>	87.24	5.44	<u>70.79</u>	75.66	50.63
	CAA	<u>91.48</u>	98.75	<u>95.12</u>	5.42	70.77	<u>75.21</u>	<u>50.47</u>
	SAE _{AXBENCH}	90.74	98.42	94.58	<u>5.43</u>	70.89	72.63	49.65
	STA (Ours)	95.78	99.33	97.56	<u>5.43</u>	70.27	71.65	49.12
Llama-3.1-8B	Vanilla	59.78	58.38	59.08	4.04	<u>58.10</u>	43.97	35.37
	Prompt _{hand}	63.70	57.30	60.50	3.62	<u>58.10</u>	46.78	36.17
	Prompt _{auto}	61.63	60.64	61.14	4.03	58.10	41.55	34.56
	CAA	68.67	<u>72.81</u>	70.74	3.89	57.64	<u>44.35</u>	<u>35.29</u>
	SAE _{AXBENCH}	<u>70.74</u>	71.98	<u>71.36</u>	<u>3.96</u>	57.88	43.44	35.09
	STA (Ours)	70.81	73.64	72.23	3.92	58.29	39.35	33.85


在Gemma和LLaMA系列模型上的实验表明，STA方法能够有效抑制越狱攻击带来的违规输出，同时不削弱模型在正常问题回答中的表现

大模型行为定向编辑干预方法: STA




调控 (Steering) 向量比提示工程对模型行为的影响**更鲁棒**

大模型行为定向编辑干预方法: STA



(a) The boundary of steering.




(b) The boundary of prompting.

调控 (Steering) 向量具有相对较好的灵活性


大模型行为定向编辑干预方法: RICE

- DeepSeek-R1 存在**语义专家路由**现象
- 仅需两个专家即可实现思维调控，而无需额外训练

DeepSeek-R1: **671B**



(a) Cognitive Experts



(b) Reasoning Accuracy

创新点：利用 nPMI 指标精准识别并放大“认知专家”的作用，从而在推理任务中实现思维过程增强，揭示了**专家（MoE）层面干预提升大模型推理能力**的可行路径

Two Experts Are All You Need for Steering Thinking:

Reinforcing Cognitive Effort in MoE Reasoning Models Without Additional Training (NeurIPS 2025)

更新: 知识 (记忆) 学习与更新

符号知识 (流程) 迁移




Figure 1: With **procedural memory**, agents can improve both the success rate (accuracy ↑) and execution efficiency (steps ↓) when solving similar tasks.

流程知识 (记忆)

符号知识 (流程) 迁移




Figure 2: The procedural memory framework consists of **Build**, **Retrieve**, and **Update**, which respectively involve encoding stored procedural memory, forming new procedural memories, and modifying existing ones in light of new experiences.

流程知识（记忆）

符号知识（流程）迁移

Model	Granularity	TravelPlanner			ALFWorld		
		#CS ↑	#HC ↑	Steps ↓	Dev ↑	Test ↑	Steps ↓
GPT-4o	No Memory	71.93	12.88	17.84	39.28	42.14	23.76
	Script	72.08	5.50	15.79	66.67	56.43	18.52
	Trajectory	<u>76.02</u>	8.25	<u>14.64</u>	67.17	74.29	<u>16.49</u>
	Proceduralization	79.94	<u>9.76</u>	14.62	87.14	77.86	15.01
Claude-3.5-sonnet	No Memory	63.49	33.06	18.84	39.20	34.97	24.12
	Script	62.08	29.61	19.21	56.13	53.59	19.38
	Trajectory	<u>65.76</u>	29.61	<u>17.72</u>	69.28	<u>71.78</u>	<u>15.97</u>
	Proceduralization	65.46	<u>30.14</u>	15.29	82.50	74.72	15.79
Qwen2.5-72b	No Memory	56.57	7.34	18.32	44.91	41.25	21.38
	Script	58.59	7.34	18.53	<u>66.24</u>	61.88	17.13
	Trajectory	<u>63.41</u>	<u>12.66</u>	<u>18.12</u>	64.49	<u>69.57</u>	<u>16.40</u>
	Proceduralization	63.82	14.19	17.94	85.71	77.19	15.32

Table 1: Results on **Build Policy**. #CS, #HC denote Commensense and Hard Constraint, respectively. ↑ indicates the higher values are better, and ↓ denotes the lower values are better. The best results among all methods with similar settings are **bolded**, and the second-best results are underlined.

Model	Policy	#CS ↑	#HC ↑	Steps ↓
GPT-4o	No Memory	71.93	12.88	17.84
	Random Sample	<u>74.59</u>	6.72	<u>15.12</u>
	Key=Query	73.38	<u>8.95</u>	15.44
	Key=AveFact	76.02	8.25	14.64
Claude-3.5-sonnet	No Memory	63.49	33.06	18.84
	Random Sample	63.99	<u>29.91</u>	17.93
	Key=Query	<u>64.93</u>	28.56	17.60
	Key=AveFact	65.76	29.61	<u>17.72</u>
Qwen2.5-72b	No Memory	56.57	7.34	18.32
	Random Sample	59.76	8.43	<u>18.31</u>
	Key=Query	<u>61.71</u>	<u>11.97</u>	18.54
	Key=AveFact	63.41	12.66	18.12

Table 2: Results on **Retrieve Policy** on TravelPlanner.

知识驱动智能体（流程记忆）持续进化大模型能力

符号知识（流程）迁移







Figure 3: Reward gain and steps reduction vs. trajectory group index with **procedural memory**.

大模型知识机理、编辑与记忆技术体系

神经网络的可塑性



可靠的开源工具栈: EasyEdit1/2系列、LightMem



可靠的评估方式: 基于开放生成和**LLM as a judge**

The Mirage of Model Editing: Revisiting Evaluation in the Wild (ACL 2025)

EasyEdit2: An Easy-to-use Steering Framework for Editing Large Language Models (EMNLP 2025 Demo)

理解与控制人工智能



图片来源于互联网

机器学习的进步最终可能会催生一个**全新的研究方向**，就像DNA结构之于生物学一样，为人工智能领域带来根本性的突破

谢谢!